

# Frameworks Comparison: Credit Scoring

## Frameworks Comparison: The Data Analyst's Toolkit

### Python (SQLAlchemy + Pandas) vs Java (JPA/JDBC + DFlib)

**Phase:** Études — Exploration des Frameworks **Contexte:** Banking Credit Scoring System

---

## Table des Matières

1. [Vue d'ensemble architecturale](#)
  2. [Le Schéma Commun — Credit Scoring Model](#)
  3. [Couche ORM: Définition des Modèles](#)
  4. [Couche Query: Écrire des Requêtes](#)
  5. [Couche DataFrame: Analyse de Données](#)
  6. [Le Workflow Complet du Data Analyst](#)
  7. [Migrations: Alembic vs Flyway](#)
  8. [Environnement Docker](#)
- 

## 1. Vue d'ensemble architecturale

Le workflow d'un data analyst suit toujours le même schéma, quel que soit le langage :

[Base de données] → [Couche ORM/Query] → [DataFrame] → [Analyse/Visualisation]

Voici comment chaque brique se mappe entre les deux écosystèmes :

Rôle	Python	Java
<b>ORM complet</b>	SQLAlchemy (Core + ORM)	JPA (Jakarta Persistence API)
<b>Query builder</b>	SQLAlchemy Core <code>select()</code>	JDBC + SQL natif
<b>Driver bas niveau</b>	<code>pymysql / asyncmy</code>	JDBC ( <code>mysql-connector-java</code> )
<b>DataFrame</b>	Pandas ( <code>pd.DataFrame</code> )	DFlib (DataFrame)

Rôle	Python	Java
<b>SQL → DataFrame</b>	pd.read_sql()	Jdbc.connector().read()
<b>Migration</b>	Alembic	Flyway
<b>Gestion deps</b>	uv (pyproject.toml)	Maven (pom.xml)

## Différences fondamentales de philosophie

**SQLAlchemy** est un framework à deux niveaux. Le **Core** est un query builder SQL pur (on écrit du SQL de manière programmatique). L'**ORM** ajoute une couche d'abstraction objet par-dessus. En tant que data analyst, tu utilises souvent les deux : l'ORM pour définir le schéma, le Core pour écrire des requêtes analytiques complexes.

**JPA** (Jakarta Persistence API) est une **spécification**, pas une implémentation. Hibernate est l'implémentation la plus répandue. JPA est purement ORM : on pense en objets, on écrit en JPQL (Java Persistence Query Language). Pour les requêtes complexes, on descend vers du **Native SQL** ou on utilise le **Criteria API**.

**JDBC** est le niveau le plus bas en Java — l'équivalent direct de `pymysql`. On écrit du SQL brut, on récupère des `ResultSet`. C'est ce que DFlib utilise directement pour charger les données.

---

## 2. Le Schéma Commun – Credit Scoring Model

Pour toutes les comparaisons, nous utilisons ce schéma relationnel bancaire :

```

clients (id PK, first_name, last_name, email, annual_income
DECIMAL,
        employment_status ENUM, credit_history_years INT,
registration_date)
    |
    < credit_applications (id PK, client_id FK, product_id
FK,
        | requested_amount DECIMAL, credit_score INT,
risk_level ENUM,
        | status ENUM, interest_rate DECIMAL,
application_date)
    |
loan_products (id PK, code, name, max_amount DECIMAL,
                base_rate DECIMAL, term_months INT)

```

Relations : - Un **client** dépose plusieurs **demandes de crédit** (`credit_applications`) - Un **produit de prêt** (`loan_product`) est référencé dans plusieurs **demandes** - `credit_applications` est la table de jonction enrichie

— elle porte le `credit_score`, le `risk_level`, le `status`, et le `interest_rate` calculé

## Contexte métier

Le **credit scoring** est le processus par lequel une banque évalue le risque d'un emprunteur. Le score (typiquement 300-850, calqué sur FICO) détermine : - Si la demande est **approuvée**, **rejetée**, ou **en révision** - Le **niveau de risque** (LOW, MEDIUM, HIGH, CRITICAL) - Le **taux d'intérêt** appliqué (plus le risque est élevé, plus le taux monte)

En tant que data analyst bancaire, les questions typiques sont : - Quel est le taux d'approbation par produit ? - Quelle est la distribution des scores par tranche de revenu ? - Quel est le taux moyen accordé par niveau de risque ? - Quels clients à haut revenu ont un score étonnamment bas ? (anomalies)

---

## 3. Couche ORM : Définition des Modèles

### 3.1 — SQLAlchemy 2.0 (Nouvelle API avec `Mapped` et `mapped_column`)

SQLAlchemy 2.0 a introduit un changement majeur : le passage de l'ancienne syntaxe `Column()` vers le système de type annotations avec `Mapped[]` et `mapped_column()`. Ce nouveau style apporte la **type safety** native — Pydantic et les IDE comprennent les types directement.

Points clés de la nouvelle API : - `DeclarativeBase` remplace `declarative_base()` - `Mapped[type]` déclare le type Python de la colonne - `mapped_column()` remplace `Column()` et infère le type SQL depuis le type Python - `relationship()` reste identique mais profite du typage `Mapped[list[...]]` - Les `Enum` Python se mappent directement vers des `ENUM MySQL`

```
# models.py – SQLAlchemy 2.0 (Sync) – Banking Credit Scoring
import enum
from datetime import date
from decimal import Decimal
from sqlalchemy import ForeignKey, String, DECIMAL, Enum,
Integer, create_engine
from sqlalchemy.orm import (
    DeclarativeBase, Mapped, mapped_column,
    relationship, sessionmaker
)

# --- Enums métier ---
class EmploymentStatus(enum.Enum):
    """Statut d'emploi du client.
    Impacte directement le scoring : un CDI est plus favorable
```

qu'un indépendant.

```
"""
SALARIED = "SALARIED"          # CDI / salarié
SELF_EMPLOYED = "SELF_EMPLOYED"  # Indépendant / freelance
UNEMPLOYED = "UNEMPLOYED"       # Sans emploi
RETIRED = "RETIRED"             # Retraité
STUDENT = "STUDENT"             # Étudiant

class RiskLevel(enum.Enum):
    """Niveau de risque calculé à partir du credit_score.
    LOW: score >= 750 | MEDIUM: 650-749 | HIGH: 550-649 |
CRITICAL: < 550
"""

    LOW = "LOW"
    MEDIUM = "MEDIUM"
    HIGH = "HIGH"
    CRITICAL = "CRITICAL"

class ApplicationStatus(enum.Enum):
    """Statut du dossier de demande de crédit."""
    PENDING = "PENDING"          # En attente d'analyse
    APPROVED = "APPROVED"        # Approuvé
    REJECTED = "REJECTED"        # Rejeté
    UNDER_REVIEW = "UNDER_REVIEW" # En révision manuelle

class Base(DeclarativeBase):
    """Classe de base pour tous les modèles.
    Remplace l'ancien declarative_base().
    """
    pass

class Client(Base):
    __tablename__ = "clients"

# Mapped[int] → INTEGER, mapped_column(primary_key=True) → PK
# auto-increment
    id: Mapped[int] = mapped_column(primary_key=True,
autoincrement=True)
        first_name: Mapped[str] = mapped_column(String(100))
        last_name: Mapped[str] = mapped_column(String(100))
        email: Mapped[str] = mapped_column(String(255), unique=True)
        annual_income: Mapped[Decimal] = mapped_column(DECIMAL(12,
2))
```

```

employment_status: Mapped[EmploymentStatus] = mapped_column(
    Enum(EmploymentStatus, native_enum=True)
)
credit_history_years: Mapped[int] = mapped_column(Integer,
default=0)
registration_date: Mapped[date] =
mapped_column(default=date.today)

# Relationship: un client a plusieurs demandes de crédit
# back_populates crée la liaison bidirectionnelle
applications: Mapped[list["CreditApplication"]] =
relationship(
    back_populates="client"
)

def __repr__(self) -> str:
    return f"Client(id={self.id}, name={self.first_name}"
{self.last_name}, income={self.annual_income})"

class LoanProduct(Base):
    __tablename__ = "loan_products"

    id: Mapped[int] = mapped_column(primary_key=True,
autoincrement=True)
    code: Mapped[str] = mapped_column(String(20), unique=True)
    name: Mapped[str] = mapped_column(String(200))
    max_amount: Mapped[Decimal] = mapped_column(DECIMAL(14, 2))
    base_rate: Mapped[Decimal] = mapped_column(DECIMAL(5, 2)) #
taux de base en %
    term_months: Mapped[int] = mapped_column(Integer)

    applications: Mapped[list["CreditApplication"]] =
relationship(
    back_populates="product"
)

def __repr__(self) -> str:
    return f"LoanProduct(id={self.id}, code={self.code},
rate={self.base_rate}%)"

class CreditApplication(Base):
    __tablename__ = "credit_applications"

    id: Mapped[int] = mapped_column(primary_key=True,
autoincrement=True)
    client_id: Mapped[int] =

```

```

mapped_column(ForeignKey("clients.id"))
    product_id: Mapped[int] =
mapped_column(ForeignKey("loan_products.id"))

        # Données de la demande
        requested_amount: Mapped[Decimal] =
mapped_column(DECIMAL(14, 2))
            credit_score: Mapped[int] = mapped_column(Integer)
# 300-850 (échelle FICO)
            risk_level: Mapped[RiskLevel]
= mapped_column(Enum(RiskLevel, native_enum=True))
            status: Mapped[ApplicationStatus] = mapped_column(
                Enum(ApplicationStatus, native_enum=True),
                default=ApplicationStatus.PENDING
            )
            interest_rate: Mapped[Decimal] = mapped_column(DECIMAL(5,
2)) # taux accordé
            application_date: Mapped[date] =
mapped_column(default=date.today)

        # Relations inverses
        client: Mapped["Client"] =
relationship(back_populates="applications")
            product: Mapped["LoanProduct"] =
relationship(back_populates="applications")

    def __repr__(self) -> str:
        return (
            f"CreditApplication(client_id={self.client_id}, "
            f"score={self.credit_score},
risk={self.risk_level.value}, "
            f"status={self.status.value})"
        )

# --- Engine et Session Factory ---
DATABASE_URL = "mysql+pymysql://analyst:analyst@mysql:3306/
credit_scoring_db"
engine = create_engine(DATABASE_URL, echo=False)
SessionLocal = sessionmaker(bind=engine)

```

## 3.2 – SQLAlchemy 2.0 (Async)

La version asynchrone change uniquement la couche connexion. Les modèles restent identiques. Ce qui change : `create_async_engine`, `async_sessionmaker`, et `AsyncSession`.

```
# models_async.py – SQLAlchemy 2.0 (Async) – Banking Credit
Scoring
```

```

from sqlalchemy.ext.asyncio import (
    create_async_engine,
    async_sessionmaker,
    AsyncSession
)

# Le driver change : pymysql → asyncmy (ou aiomysql)
ASYNC_DATABASE_URL = "mysql+asyncmy://analyst:analyst@mysql:3306/
credit_scoring_db"

async_engine = create_async_engine(ASYNC_DATABASE_URL,
echo=False)
AsyncSessionLocal = async_sessionmaker(
    bind=async_engine,
    class_=AsyncSession,
    expire_on_commit=False # Important: évite le lazy-loading
implicite en async
)

# Les modèles (Client, LoanProduct, CreditApplication) restent
EXACTEMENT les mêmes.
# Seule la manière de créer les tables et d'exécuter les requêtes
change.

# Création des tables en async :
async def init_db():
    async with async_engine.begin() as conn:
        await conn.run_sync(Base.metadata.create_all)

```

### 3.3 – JPA (Jakarta Persistence avec Hibernate)

JPA utilise des **annotations** sur les classes Java. La philosophie est similaire à SQLAlchemy ORM, mais plus rigide : tout passe par des entités, le schéma est décrit exclusivement via les annotations.

Points clés : - `@Entity` marque la classe comme table - `@Id` + `@GeneratedValue` = clé primaire auto-incrémentée - `@ManyToOne` / `@OneToOne` = relations avec lazy/eager loading - `@Enumerated(EnumType.STRING)` pour mapper les enums Java vers des ENUM SQL - `persistence.xml` configure la connexion (équivalent de `create_engine`)

```

// --- Enums métier ---

// EmploymentStatus.java
package com.creditscoring.model;

public enum EmploymentStatus {
    SALARIED,

```

```
    SELF_EMPLOYED,
    UNEMPLOYED,
    RETIRED,
    STUDENT
}

// RiskLevel.java
package com.creditscoring.model;

public enum RiskLevel {
    LOW,           // score >= 750
    MEDIUM,        // 650-749
    HIGH,          // 550-649
    CRITICAL      // < 550
}

// ApplicationStatus.java
package com.creditscoring.model;

public enum ApplicationStatus {
    PENDING,
    APPROVED,
    REJECTED,
    UNDER_REVIEW
}

// Client.java – JPA Entity
package com.creditscoring.model;

import jakarta.persistence.*;
import java.math.BigDecimal;
import java.time.LocalDate;
import java.util.ArrayList;
import java.util.List;

@Entity
@Table(name = "clients")
public class Client {

    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;

    @Column(name = "first_name", length = 100, nullable = false)
    private String firstName;

    @Column(name = "last_name", length = 100, nullable = false)
    private String lastName;
```

```

    @Column(length = 255, unique = true, nullable = false)
    private String email;

    @Column(name = "annual_income", precision = 12, scale = 2,
nullable = false)
    private BigDecimal annualIncome;

    // @Enumerated(EnumType.STRING) stocke le NOM de l'enum
("SALARIED")
    // EnumType.ORDINAL stockerait l'index (0, 1, 2) – à éviter
absolument
    @Enumerated(EnumType.STRING)
    @Column(name = "employment_status", nullable = false)
    private EmploymentStatus employmentStatus;

    @Column(name = "credit_history_years", nullable = false)
    private Integer creditHistoryYears = 0;

    @Column(name = "registration_date")
    private LocalDate registrationDate = LocalDate.now();

    // OneToMany: un client a plusieurs demandes de crédit
    // mappedBy pointe vers le champ "client" dans
CreditApplication
    @OneToMany(mappedBy = "client", cascade = CascadeType.ALL,
fetch = FetchType.LAZY)
    private List<CreditApplication> applications = new
ArrayList<>();

    // Constructeur vide obligatoire pour JPA
    public Client() {}

    public Client(String firstName, String lastName, String
email,
                           BigDecimal annualIncome, EmploymentStatus
employmentStatus,
                           int creditHistoryYears) {
        this.firstName = firstName;
        this.lastName = lastName;
        this.email = email;
        this.annualIncome = annualIncome;
        this.employmentStatus = employmentStatus;
        this.creditHistoryYears = creditHistoryYears;
    }

    // Getters
    public Long getId() { return id; }

```

```
    public String getFirstName() { return firstName; }
    public String getLastName() { return lastName; }
    public String getEmail() { return email; }
    public BigDecimal getAnnualIncome() { return annualIncome; }
    public EmploymentStatus getEmploymentStatus() { return
employmentStatus; }
        public Integer getCreditHistoryYears() { return
creditHistoryYears; }
        public LocalDate getRegistrationDate() { return
registrationDate; }
        public List<CreditApplication> getApplications() { return
applications; }
    }

// LoanProduct.java – JPA Entity
package com.creditscoring.model;

import jakarta.persistence.*;
import java.math.BigDecimal;
import java.util.ArrayList;
import java.util.List;

@Entity
@Table(name = "loan_products")
public class LoanProduct {

    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;

    @Column(length = 20, unique = true, nullable = false)
    private String code;

    @Column(length = 200, nullable = false)
    private String name;

    @Column(name = "max_amount", precision = 14, scale = 2,
nullable = false)
    private BigDecimal maxAmount;

    @Column(name = "base_rate", precision = 5, scale = 2,
nullable = false)
    private BigDecimal baseRate;

    @Column(name = "term_months", nullable = false)
    private Integer termMonths;

    @OneToMany(mappedBy = "product", cascade = CascadeType.ALL,

```

```
fetch = FetchType.LAZY)
    private List<CreditApplication> applications = new
ArrayList<>();

    public LoanProduct() {}

    public LoanProduct(String code, String name, BigDecimal
maxAmount,
                           BigDecimal baseRate, int termMonths) {
        this.code = code;
        this.name = name;
        this.maxAmount = maxAmount;
        this.baseRate = baseRate;
        this.termMonths = termMonths;
    }

    public Long getId() { return id; }
    public String getCode() { return code; }
    public String getName() { return name; }
    public BigDecimal getMaxAmount() { return maxAmount; }
    public BigDecimal getBaseRate() { return baseRate; }
    public Integer getTermMonths() { return termMonths; }
}

// CreditApplication.java – JPA Entity
package com.creditscoring.model;

import jakarta.persistence.*;
import java.math.BigDecimal;
import java.time.LocalDate;

@Entity
@Table(name = "credit_applications")
public class CreditApplication {

    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;

    // ManyToOne: plusieurs demandes appartiennent à un client
    @ManyToOne(fetch = FetchType.LAZY)
    @JoinColumn(name = "client_id", nullable = false)
    private Client client;

    @ManyToOne(fetch = FetchType.LAZY)
    @JoinColumn(name = "product_id", nullable = false)
    private LoanProduct product;
```

```
        @Column(name = "requested_amount", precision = 14, scale = 2, nullable = false)
    private BigDecimal requestedAmount;

        @Column(name = "credit_score", nullable = false)
    private Integer creditScore; // 300-850

        @Enumerated(EnumType.STRING)
    @Column(name = "risk_level", nullable = false)
    private RiskLevel riskLevel;

        @Enumerated(EnumType.STRING)
    @Column(nullable = false)
    private ApplicationStatus status = ApplicationStatus.PENDING;

        @Column(name = "interest_rate", precision = 5, scale = 2, nullable = false)
    private BigDecimal interestRate;

        @Column(name = "application_date")
    private LocalDate applicationDate = LocalDate.now();

    public CreditApplication() {}

    public CreditApplication(Client client, LoanProduct product,
                           BigDecimal requestedAmount, int creditScore,
                           RiskLevel riskLevel, BigDecimal interestRate) {
        this.client = client;
        this.product = product;
        this.requestedAmount = requestedAmount;
        this.creditScore = creditScore;
        this.riskLevel = riskLevel;
        this.interestRate = interestRate;
    }

    public Long getId() { return id; }
    public Client getClient() { return client; }
    public LoanProduct getProduct() { return product; }
    public BigDecimal getRequestedAmount() { return requestedAmount; }
    public Integer getCreditScore() { return creditScore; }
    public RiskLevel getRiskLevel() { return riskLevel; }
    public ApplicationStatus getStatus() { return status; }
    public BigDecimal getInterestRate() { return interestRate; }
    public LocalDate getApplicationDate() { return
```

```
applicationDate; }
    }
```

### 3.4 – Comparaison directe : Mapping des concepts ORM

Concept	SQLAlchemy 2.0	JPA / Hibernate
Déclarer un modèle	class Client(Base):	@Entity class Client {}
Clé primaire	Mapped[int] = mapped_column(primary_key=True)	@Id @GeneratedValue
Colonne typée	Mapped[str] = mapped_column(String(100))	@Column(length = 100)
Colonne décimale	Mapped[Decimal] = mapped_column(DECIMAL(12,2))	@Column(precision=12, scale=2)
Enum	Mapped[RiskLevel] = mapped_column(Enum(RiskLevel))	@Enumerated(EnumType.STRING)
Clé étrangère	mapped_column(ForeignKey("clients.id"))	@JoinColumn(name = "client_id")
Relation 1-N	Mapped[list["CreditApplication"]] = relationship()	@OneToMany(mappedBy = "client")
Relation N-1	Mapped["Client"] = relationship()	@ManyToOne @JoinColumn
Lazy loading	Défaut pour les relations	fetch = FetchType.LAZY
Session / Transaction	Session() / with session:	EntityManager / @Transactional
Config connexion	create_engine(url)	persistence.xml

## 4. Couche Query : Écrire des Requêtes

C'est ici que la différence est la plus significative pour un data analyst bancaire. Tu écris des requêtes analytiques (JOINS, GROUP BY, agrégations, CASE WHEN) et tu veux récupérer des données tabulaires.

### 4.1 – SQLAlchemy 2.0 Core (Nouvelle API `select()`)

La nouvelle API abandonne l'ancien style `session.query(Model)` au profit de `select()` combiné avec `session.execute()`. C'est plus explicite et plus composable.

```
# queries_sync.py - SQLAlchemy 2.0 Select API - Credit Scoring
from sqlalchemy import select, func, desc, case, and_, cast,
Float

# ----- QUERY 1: Toutes les demandes avec détails client +
produit -----
```

```

# Ancien style (déprécié) :
session.query(CreditApplication).all()
# Nouveau style :
stmt_all = (
    select(
        Client.first_name,
        Client.last_name,
        Client.annual_income,
        Client.employment_status,
        LoanProduct.name.label("product_name"),
        LoanProduct.base_rate,
        CreditApplication.requested_amount,
        CreditApplication.credit_score,
        CreditApplication.risk_level,
        CreditApplication.status,
        CreditApplication.interest_rate,
        CreditApplication.application_date
    )
    .join(CreditApplication.client)
    .join(CreditApplication.product)
    .order_by(desc(CreditApplication.application_date))
)

with SessionLocal() as session:
    result = session.execute(stmt_all)
    rows = result.all() # Liste de Row tuples (named tuples)
    # Chaque row: row.first_name, row.credit_score,
    row.risk_level, etc.

# ----- QUERY 2: Taux d'approbation par produit -----
# SELECT p.name, COUNT(*), SUM(CASE WHEN status='APPROVED' THEN 1
ELSE 0 END),
#       SUM(...) / COUNT(*) * 100 AS approval_rate
stmt_approval = (
    select(
        LoanProduct.code,
        LoanProduct.name,
        func.count(CreditApplication.id).label("total_applications"),
        func.sum(
            case(
                (CreditApplication.status ==
ApplicationStatus.APPROVED, 1),
                else_=0
            )
        ).label("approved_count"),
        (

```

```

        func.sum(
            case(
                (CreditApplication.status ==
ApplicationStatus.APPROVED, 1),
                else_=0
            )
        ) * 100.0 / func.count(CreditApplication.id)
    ).label("approval_rate_pct")
)
.join(CreditApplication.product)
.group_by(LoanProduct.id)
.order_by(desc("approval_rate_pct"))
)

# ----- QUERY 3: Statistiques par niveau de risque -----
stmt_risk = (
    select(
        CreditApplication.risk_level,
        func.count(CreditApplication.id).label("num_applications"),
        func.avg(CreditApplication.credit_score).label("avg_score"),
        func.avg(CreditApplication.interest_rate).label("avg_rate"),
        func.avg(CreditApplication.requested_amount).label("avg_amount"),
        func.min(CreditApplication.credit_score).label("min_score"),
        func.max(CreditApplication.credit_score).label("max_score")
    )
    .group_by(CreditApplication.risk_level)
    .order_by(CreditApplication.risk_level)
)

# ----- QUERY 4: Clients haut revenu avec score bas (anomalies) -----
# Subquery: moyenne globale des scores
global_avg_score =
select(func.avg(CreditApplication.credit_score)).scalar_subquery()

stmt_anomalies = (
    select(
        Client.first_name,
        Client.last_name,

```

```

        Client.annual_income,
        Client.employment_status,

func.avg(CreditApplication.credit_score).label("avg_score")
)
.join(CreditApplication.client)
.group_by(Client.id)
.having(
    and_(
        Client.annual_income > 100000,
# Haut revenu
        func.avg(CreditApplication.credit_score) <
global_avg_score # Score sous la moyenne
    )
)
.order_by(Client.annual_income.desc())
)

-----  

# ----- QUERY 5: Montant total de crédit exposé par statut
-----  

stmt_exposure = (
    select(
        CreditApplication.status,
        func.count(CreditApplication.id).label("count"),

func.sum(CreditApplication.requested_amount).label("total_exposure"),

func.avg(CreditApplication.requested_amount).label("avg_amount")
)
.group_by(CreditApplication.status)
)

```

## 4.2 – SQLAlchemy 2.0 Async

```

# queries_async.py – SQLAlchemy 2.0 Async – Credit Scoring
import asyncio
from sqlalchemy import select, func, desc
from sqlalchemy.orm import selectinload

# Même statements que le sync – seule l'exécution change
async def get_risk_stats():
    stmt = (
        select(
            CreditApplication.risk_level,
            func.count(CreditApplication.id).label("num_applications"),

```

```

func.avg(CreditApplication.credit_score).label("avg_score"),
func.avg(CreditApplication.interest_rate).label("avg_rate")
    )
    .group_by(CreditApplication.risk_level)
)

async with AsyncSessionLocal() as session:
    result = await session.execute(stmt)
    return result.all()

# Pour les relations (lazy loading interdit en async) :
async def get_client_with_applications(client_id: int):
    stmt = (
        select(Client)
        .options(selectinload(Client.applications)) # Eager
loading explicite
        .where(Client.id == client_id)
    )
    async with AsyncSessionLocal() as session:
        result = await session.execute(stmt)
        return result.scalar_one_or_none()

# Exécution
# asyncio.run(get_risk_stats())

```

**Point critique `async`** : En mode asynchrone, le **lazy loading implicite** est **interdit**. Si tu accèdes à `client.applications` sans avoir fait un eager load, SQLAlchemy lève une `MissingGreenlet` exception. Tu dois toujours utiliser `selectinload()`, `joinedload()`, ou `subqueryload()` explicitement.

### 4.3 – JPA (JPQL et Criteria API)

JPQL est le langage de requête de JPA. Il ressemble à SQL mais opère sur les **entités** (noms de classes) plutôt que sur les **tables**.

```

// queries_jpa.java – JPA JPQL Queries – Credit Scoring
import jakarta.persistence.*;
import java.util.List;

EntityManagerFactory emf =
Persistence.createEntityManagerFactory("credit-scoring-pu");
EntityManager em = emf.createEntityManager();

// ----- QUERY 1: Toutes les demandes avec détails
-----
TypedQuery<Object[]> q1 = em.createQuery(
    """
        SELECT
            c.id AS id,
            c.name AS name,
            c.type AS type,
            c.value AS value
        FROM
            CreditApplication c
        WHERE
            c.id = ?
    """
);
q1.setParameter(1, id);
List<Object[]> results = q1.getResultList();

```

```

        SELECT c.firstName, c.lastName, c.annualIncome,
               p.name, ca.requestedAmount, ca.creditScore,
               ca.riskLevel, ca.status, ca.interestRate
        FROM CreditApplication ca
        JOIN ca.client c
        JOIN ca.product p
        ORDER BY ca.applicationDate DESC
        """", Object[].class
    );
List<Object[]> allApps = q1.getResultList();

// ----- QUERY 2: Taux d'approbation par produit -----
TypedQuery<Object[]> q2 = em.createQuery(
    """
        SELECT p.code, p.name,
               COUNT(ca),
               SUM(CASE WHEN ca.status = 'APPROVED' THEN 1 ELSE 0
END),
               SUM(CASE WHEN ca.status = 'APPROVED' THEN 1.0 ELSE 0.0
END)
               / COUNT(ca) * 100
        FROM CreditApplication ca
        JOIN ca.product p
        GROUP BY p.id
        ORDER BY SUM(CASE WHEN ca.status = 'APPROVED' THEN 1.0 ELSE
0.0 END)
               / COUNT(ca) DESC
    """", Object[].class
);

// ----- QUERY 3: Stats par niveau de risque -----
TypedQuery<Object[]> q3 = em.createQuery(
    """
        SELECT ca.riskLevel,
               COUNT(ca),
               AVG(ca.creditScore),
               AVG(ca.interestRate),
               AVG(ca.requestedAmount),
               MIN(ca.creditScore),
               MAX(ca.creditScore)
        FROM CreditApplication ca
        GROUP BY ca.riskLevel
    """", Object[].class
);

```

```

// ----- QUERY 4: Anomalies – haut revenu / score bas
-----
TypedQuery<Object[]> q4 = em.createQuery(
    """
        SELECT c.firstName, c.lastName, c.annualIncome,
        AVG(ca.creditScore)
        FROM CreditApplication ca
        JOIN ca.client c
        GROUP BY c.id
        HAVING c.annualIncome > 100000
            AND AVG(ca.creditScore) < (SELECT AVG(ca2.creditScore)
        FROM CreditApplication ca2)
        ORDER BY c.annualIncome DESC
    """, Object[].class
);

em.close();
emf.close();

```

## 4.4 – JDBC (SQL Natif)

JDBC est le plus bas niveau. Tu écris du SQL pur, tu gères les connexions, tu itères sur les ResultSet. C'est ce que DFlib utilise en interne.

```

// queries_jdbc.java – JDBC Raw SQL – Credit Scoring
import java.sql.*;

String url = "jdbc:mysql://mysql:3306/credit_scoring_db";
String user = "analyst";
String password = "analyst";

// ----- QUERY 1: Toutes les demandes -----
String sqlAll = """
    SELECT c.first_name, c.last_name, c.annual_income,
    c.employment_status,
        p.name AS product_name, p.base_rate,
        ca.requested_amount, ca.credit_score, ca.risk_level,
        ca.status, ca.interest_rate, ca.application_date
    FROM credit_applications ca
    JOIN clients c ON ca.client_id = c.id
    JOIN loan_products p ON ca.product_id = p.id
    ORDER BY ca.application_date DESC
""";

try (Connection conn = DriverManager.getConnection(url, user,
password);
PreparedStatement pstmt = conn.prepareStatement(sqlAll));

```

```

        ResultSet rs = pstmt.executeQuery() {
            while (rs.next()) {
                String name = rs.getString("first_name") + " " +
rs.getString("last_name");
                int score = rs.getInt("credit_score");
                String risk = rs.getString("risk_level");
                double rate = rs.getDouble("interest_rate");
            }
        }

// ----- QUERY 2: Taux d'approbation par produit -----
String sqlApproval = """
    SELECT p.code, p.name,
           COUNT(*) AS total,
           SUM(CASE WHEN ca.status = 'APPROVED' THEN 1 ELSE 0
END) AS approved,
           SUM(CASE WHEN ca.status = 'APPROVED' THEN 1 ELSE 0
END) * 100.0
           / COUNT(*) AS approval_rate
    FROM credit_applications ca
    JOIN loan_products p ON ca.product_id = p.id
    GROUP BY p.id
    ORDER BY approval_rate DESC
""";
// Même pattern: prepareStatement → executeQuery → iterate
ResultSet

```

## 4.5 – Comparaison directe : Style de requêtes

Aspect	<b>SQLAlchemy 2.0 Core</b>	<b>JPQL (JPA)</b>	<b>JDBC</b>
Langage	Python DSL ( <code>select()</code> )	String JPQL	String SQL natif
JOIN syntax	<code>.join(CreditApplication.client)</code>	<code>ca.client</code>	<code>clients c</code> <code>JOIN ... ON ...</code>
CASE WHEN	<code>case((condition, val), else_=)</code>	<code>CASE WHEN ... THEN ... END</code>	SQL standard
Enum handling	Comparaison Python native	String comparison en JPQL	Comparaison en SQL
Noms utilisés	Classes + attributs Python	Classes + champs Java	Tables + colonnes SQL
Type safety	Oui (IDE autocomplete)		

Aspect	SQLAlchemy 2.0 Core	JPQL (JPA)	JDBC
Résultat	Row named tuples	Non (strings)	Non (strings)
Composabilité	Excellent (chaîner)	Object[]	ResultSet
CTE / Window funcs	Supporté nativement	Limitée Non en JPQL, Native Query	Aucune SQL complet

## 5. Couche DataFrame : Analyse de Données

C'est le cœur du workflow data analyst bancaire. On a nos données de scoring, on veut les analyser en mémoire.

### 5.1 – Pandas : SQL → DataFrame → Analyse

```
# analysis_pandas.py - Pandas Workflow - Credit Scoring
import pandas as pd
from sqlalchemy import create_engine

engine = create_engine("mysql+pymysql://analyst:analyst@mysql:
3306/credit_scoring_db")

# ===== CHARGEMENT DES DONNÉES =====

df = pd.read_sql(
    """
    SELECT c.first_name, c.last_name, c.annual_income,
           c.employment_status, c.credit_history_years,
           p.code AS product_code, p.name AS product_name,
           p.base_rate, p.term_months,
           ca.requested_amount, ca.credit_score, ca.risk_level,
           ca.status, ca.interest_rate, ca.application_date
    FROM credit_applications ca
    JOIN clients c ON ca.client_id = c.id
    JOIN loan_products p ON ca.product_id = p.id
    """,
    engine,
    parse_dates=["application_date"]
)

# ===== ANALYSE =====

# --- Taux d'approbation par produit ---
```

```

approval_by_product = (
    df.groupby("product_name")
    .agg(
        total=("status", "count"),
        approved=("status", lambda x: (x == "APPROVED").sum()),
    )
)
approval_by_product["approval_rate"] = (
    (approval_by_product["approved"] /
    approval_by_product["total"] * 100).round(2)
)
approval_by_product =
approval_by_product.sort_values("approval_rate", ascending=False)

# --- Statistiques par niveau de risque ---
risk_stats = (
    df.groupby("risk_level")["credit_score"]
    .agg(["mean", "std", "min", "max", "count"])
    .rename(columns={"mean": "avg_score", "count": "num_applications"})
    .round(2)
)

# --- Distribution des scores (binning calqué sur FICO) ---
df["score_bracket"] = pd.cut(
    df["credit_score"],
    bins=[300, 550, 650, 750, 850],
    labels=["CRITICAL (300-549)", "HIGH (550-649)", "MEDIUM (650-749)", "LOW (750-850)"]
)
score_distribution =
df["score_bracket"].value_counts().sort_index()

# --- Spread: différence entre taux accordé et taux de base ---
df["rate_spread"] = df["interest_rate"] - df["base_rate"]

spread_by_risk = (
    df.groupby("risk_level")["rate_spread"]
    .agg(["mean", "min", "max"])
    .round(2)
)

# --- Exposition totale par statut ---
exposure = (

```

```

df.groupby("status")
    .agg(
        count=("requested_amount", "count"),
        total_exposure=("requested_amount", "sum"),
        avg_amount=("requested_amount", "mean")
    )
    .round(2)
)

# --- Pivot: Produit × Risque → montant moyen demandé ---
pivot = df.pivot_table(
    values="requested_amount",
    index="product_name",
    columns="risk_level",
    aggfunc="mean"
).round(2)

# --- Anomalies: haut revenu, score bas ---
global_avg = df["credit_score"].mean()
anomalies = (
    df[
        (df["annual_income"] > 100_000) &
        (df["credit_score"] < global_avg)
    ]
    [["first_name", "last_name", "annual_income",
"credit_score", "risk_level"]]
    .drop_duplicates()
    .sort_values("annual_income", ascending=False)
)

# --- Top 10 clients par score moyen ---
top_clients = (
    df.groupby(["first_name", "last_name"])
    .agg(
        avg_score=("credit_score", "mean"),
        num_applications=("status", "count"),
        total_requested=("requested_amount", "sum")
    )
    .sort_values("avg_score", ascending=False)
    .head(10)
    .round(2)
)

```

## 5.2 – DFlib : SQL → DataFrame → Analyse

DFlib est la réponse Java à Pandas. La philosophie est similaire : des DataFrames immuables avec des opérations chaînées. La différence majeure est que DFlib est **immutable** — chaque opération retourne un nouveau DataFrame.

```
// analysis_dflib.java - DFlib Workflow - Credit Scoring
import org.dflib.*;
import org.dflib.jdbc.connector.JdbcConnector;
import org.dflib.Exp;
import org.dflib.agg.Agg;

// ===== CONNEXION JDBC =====
JdbcConnector connector = new JdbcConnector(
    "jdbc:mysql://mysql:3306/credit_scoring_db",
    "analyst",
    "analyst"
);

// ===== CHARGEMENT DES DONNÉES =====

DataFrame df = connector
    .sqlLoader("""
        SELECT c.first_name, c.last_name, c.annual_income,
               c.employment_status, c.credit_history_years,
               p.code AS product_code, p.name AS product_name,
               p.base_rate, p.term_months,
               ca.requested_amount, ca.credit_score,
               ca.risk_level,
               ca.status, ca.interest_rate, ca.application_date
        FROM credit_applications ca
        JOIN clients c ON ca.client_id = c.id
        JOIN loan_products p ON ca.product_id = p.id
    """)
    .load();

System.out.println(df.head(5));
System.out.println("Shape: " + df.height() + " rows x " +
df.width() + " cols");

// ===== ANALYSE =====

// --- Statistiques par niveau de risque ---
DataFrame riskStats = df
    .group("risk_level")
    .agg(
```

```

        Exp.$int("credit_score").avg().as("avg_score"),
        Exp.$int("credit_score").min().as("min_score"),
        Exp.$int("credit_score").max().as("max_score"),
        Exp.$decimal("interest_rate").avg().as("avg_rate"),
        Exp.$decimal("requested_amount").avg().as("avg_amount"),
        Exp.count().as("num_applications")
    )
    .sort(Exp.$decimal("avg_score").desc());
}

// --- Colonne calculée: spread (taux accordé - taux de base) ---
DataFrame withSpread = df.addColumn(
    "rate_spread",
    Exp.$decimal("interest_rate").sub(Exp.$decimal("base_rate"))
);

// --- Spread moyen par risque ---
DataFrame spreadByRisk = withSpread
    .group("risk_level")
    .agg(
        Exp.$decimal("rate_spread").avg().as("avg_spread"),
        Exp.$decimal("rate_spread").min().as("min_spread"),
        Exp.$decimal("rate_spread").max().as("max_spread")
    );
}

// --- Filtrage: demandes à haut risque rejetées ---
DataFrame highRiskRejected = df
    .rows(
        Exp.$str("risk_level").eq("CRITICAL")
            .and(Exp.$str("status").eq("REJECTED"))
    )
    .select();

// --- Exposition par statut ---
DataFrame exposure = df
    .group("status")
    .agg(
        Exp.count().as("count"),
        Exp.
$decimal("requested_amount").sum().as("total_exposure"),
        Exp.$decimal("requested_amount").avg().as("avg_amount")
    );
}

// --- Top 10 clients par score moyen ---
DataFrame topClients = df

```

```

    .group("first_name", "last_name")
    .agg(
        Exp.$int("credit_score").avg().as("avg_score"),
        Exp.count().as("num_applications"),
        Exp.

$decimal("requested_amount").sum().as("total_requested")
    )
    .sort(Exp.$decimal("avg_score").desc())
    .head(10);

// --- Export CSV ---
Csv.saver().save(riskStats, "risk_analysis.csv");

```

### 5.3 – Comparaison directe : Pandas vs DFlib

Opération	Pandas	DFlib
<b>Charger depuis SQL</b>	pd.read_sql(sql, engine)	connector.sqlLoader(sql).load()
<b>Voir les premières lignes</b>	df.head(5)	df.head(5)
<b>Sélectionner des colonnes</b>	df[["col1", "col2"]]	df.cols("col1", "col2").select()
<b>Filtrer les lignes</b>	df[df["status"] == "APPROVED"]	df.rows(Exp.\$str("status").eq("APPROVED"))
<b>Filtrer combiné (AND)</b>	df[(cond1) & (cond2)]	df.rows(Exp.\$str("x").eq("A") \$int("y").gt(5)).select()
<b>Ajouter une colonne</b>	df["spread"] = df["rate"] - df["base"]	df.addColumn("spread", Exp.\$decimal("rate").sub(Exp.\$decimal("base")))
<b>Group by + agg</b>	df.groupby("risk").agg({"score": "mean"})	df.group("risk").agg(Exp.\$int("score").avg().as("mean"))
<b>Lambda / custom agg</b>	.agg(lambda x: (x == "APPROVED").sum())	Pas de lambda — SQL-side ou post
<b>Binning (pd.cut)</b>	pd.cut(df["score"], bins=[...])	Pas de cut natif — condition chaînée
<b>Pivot table</b>	df.pivot_table(values, index, columns)	Pas de pivot natif — group + reshape
<b>Trier</b>	df.sort_values("col", ascending=False)	df.sort(Exp.\$decimal("col").desc())
<b>Merge/Join</b>	pd.merge(df1, df2, on="key")	df1.join(df2).on("key").select()
<b>Mutabilité</b>	<b>Mutable</b> par défaut	<b>Immutable</b> — chaque op retourne nouveau DF
<b>Type system</b>	Dynamique (dtype inféré)	Fort (\$int(), \$decimal(), \$str())

# 6. Le Workflow Complet du Data Analyst

## Python — Le workflow que tu connais

```
# workflow_python.py - Credit Scoring Notebook
from sqlalchemy import create_engine
import pandas as pd

# 1. Connexion
engine = create_engine("mysql+pymysql://analyst:analyst@mysql:
3306/credit_scoring_db")

# 2. Query custom
sql = """
    SELECT c.first_name, c.last_name, c.annual_income,
c.employment_status,
        p.code, p.name AS product, p.base_rate, p.term_months,
        ca.requested_amount, ca.credit_score, ca.risk_level,
        ca.status, ca.interest_rate, ca.application_date
    FROM credit_applications ca
    JOIN clients c ON ca.client_id = c.id
    JOIN loan_products p ON ca.product_id = p.id
"""

# 3. Chargement en DataFrame
df = pd.read_sql(sql, engine, parse_dates=["application_date"])

# 4. Analyse rapide
print(df.describe())
print(df.dtypes)
print(df.shape)

# 5. Transformations
df["rate_spread"] = df["interest_rate"] - df["base_rate"]
df["income_bracket"] = pd.cut(
    df["annual_income"],
    bins=[0, 30000, 60000, 100000, 200000, float("inf")],
    labels=["<30K", "30-60K", "60-100K", "100-200K", "200K+"]
)

# 6. Agrégation
risk_report = (
    df.groupby("risk_level")
    .agg(
        avg_score=("credit_score", "mean"),
        avg_rate=("interest_rate", "mean"),
        total_exposure=("requested_amount", "sum"),
        count=("status", "count")
    )
)
```

```

        )
    .round(2)
)

# 7. Export
risk_report.to_csv("credit_risk_report.csv")

```

## Java – Le même workflow avec DFlib

```

// workflow_java.java – Credit Scoring Notebook (IJava kernel)
import org.dflib.*;
import org.dflib.jdbc.connector.JdbcConnector;

// 1. Connexion
JdbcConnector connector = new JdbcConnector(
    "jdbc:mysql://mysql:3306/credit_scoring_db", "analyst",
"analyst"
);

// 2 + 3. Query + Chargement en DataFrame
DataFrame df = connector.sqlLoader("""
    SELECT c.first_name, c.last_name, c.annual_income,
c.employment_status,
            p.code, p.name AS product, p.base_rate, p.term_months,
            ca.requested_amount, ca.credit_score, ca.risk_level,
            ca.status, ca.interest_rate, ca.application_date
    FROM credit_applications ca
    JOIN clients c ON ca.client_id = c.id
    JOIN loan_products p ON ca.product_id = p.id
""").load();

// 4. Analyse rapide
System.out.println(df.head(5));
System.out.println("Shape: " + df.height() + " rows x " +
df.width() + " cols");
System.out.println("Columns: " + df.getColumnsIndex());

// 5. Transformation
DataFrame enriched = df.addColumn(
    "rate_spread",
    Exp.$decimal("interest_rate").sub(Exp.$decimal("base_rate"))
);

// 6. Agrégation
DataFrame riskReport = enriched
    .group("risk_level")
    .agg(
        Exp.$int("credit_score").avg().as("avg_score"),

```

```

        Exp.$decimal("interest_rate").avg().as("avg_rate"),
        Exp.
$decimal("requested_amount").sum().as("total_exposure"),
        Exp.count().as("count")
    )
    .sort(Exp.$decimal("avg_score").desc()));

// 7. Export
Csv.saver().save(riskReport, "credit_risk_report.csv");

```

---

## 7. Migrations : Alembic vs Flyway

### 7.1 — Alembic (Python)

```

# Installation et initialisation
uv add alembic
alembic init migrations
# Configurer alembic.ini : sqlalchemy.url = mysql+pymysql://
analyst:analyst@mysql:3306/credit_scoring_db
# Configurer migrations/env.py : target_metadata = Base.metadata

# $ alembic revision --autogenerate -m "create credit scoring
tables"
# Fichier généré: migrations/versions/
xxxx_create_credit_scoring_tables.py
def upgrade():
    op.create_table('clients',
                    sa.Column('id', sa.Integer(), autoincrement=True,
nullable=False),
                    sa.Column('first_name', sa.String(100), nullable=False),
                    sa.Column('last_name', sa.String(100), nullable=False),
                    sa.Column('email', sa.String(255), nullable=False),
                    sa.Column('annual_income', sa.DECIMAL(12, 2),
nullable=False),
                    sa.Column('employment_status',
sa.Enum('SALARIED', 'SELF_EMPLOYED', 'UNEMPLOYED', 'RETIRED', 'STUDENT'),
nullable=False),
                    sa.Column('credit_history_years', sa.Integer(),
nullable=False),
                    sa.Column('registration_date', sa.Date(),
nullable=False),
                    sa.PrimaryKeyConstraint('id'),
                    sa.UniqueConstraint('email')
    )
    op.create_table('loan_products',
                    sa.Column('id', sa.Integer(), autoincrement=True,
nullable=False),

```

```

        sa.Column('code', sa.String(20), nullable=False),
        sa.Column('name', sa.String(200), nullable=False),
        sa.Column('max_amount', sa.DECIMAL(14, 2),
nullable=False),
        sa.Column('base_rate', sa.DECIMAL(5, 2), nullable=False),
        sa.Column('term_months', sa.Integer(), nullable=False),
        sa.PrimaryKeyConstraint('id'),
        sa.UniqueConstraint('code')
    )
    op.create_table('credit_applications',
        sa.Column('id', sa.Integer(), autoincrement=True,
nullable=False),
        sa.Column('client_id', sa.Integer(), nullable=False),
        sa.Column('product_id', sa.Integer(), nullable=False),
        sa.Column('requested_amount', sa.DECIMAL(14, 2),
nullable=False),
        sa.Column('credit_score', sa.Integer(), nullable=False),
        sa.Column('risk_level',
sa.Enum('LOW', 'MEDIUM', 'HIGH', 'CRITICAL'), nullable=False),
        sa.Column('status',
sa.Enum('PENDING', 'APPROVED', 'REJECTED', 'UNDER REVIEW'),
nullable=False),
        sa.Column('interest_rate', sa.DECIMAL(5, 2),
nullable=False),
        sa.Column('application_date', sa.Date(), nullable=False),
        sa.ForeignKeyConstraint(['client_id'], ['clients.id']),
        sa.ForeignKeyConstraint(['product_id'],
['loan_products.id']),
        sa.PrimaryKeyConstraint('id')
    )

    def downgrade():
        op.drop_table('credit_applications')
        op.drop_table('loan_products')
        op.drop_table('clients')

    alembic upgrade head      # Appliquer
    alembic history          # Voir l'historique
    alembic downgrade -1     # Rollback

```

## 7.2 – Flyway (Java)

```

-- V1_create_clients.sql
CREATE TABLE clients (
    id BIGINT AUTO_INCREMENT PRIMARY KEY,
    first_name VARCHAR(100) NOT NULL,
    last_name VARCHAR(100) NOT NULL,
    email VARCHAR(255) NOT NULL UNIQUE,

```

```

        annual_income DECIMAL(12,2) NOT NULL,
        employment_status
ENUM('SALARIED','SELF_EMPLOYED','UNEMPLOYED','RETIRED','STUDENT')
NOT NULL,
        credit_history_years INT NOT NULL DEFAULT 0,
        registration_date DATE NOT NULL DEFAULT (CURRENT_DATE)
);

-- V2_create_loan_products.sql
CREATE TABLE loan_products (
    id BIGINT AUTO_INCREMENT PRIMARY KEY,
    code VARCHAR(20) NOT NULL UNIQUE,
    name VARCHAR(200) NOT NULL,
    max_amount DECIMAL(14,2) NOT NULL,
    base_rate DECIMAL(5,2) NOT NULL,
    term_months INT NOT NULL
);

-- V3_create_credit_applications.sql
CREATE TABLE credit_applications (
    id BIGINT AUTO_INCREMENT PRIMARY KEY,
    client_id BIGINT NOT NULL,
    product_id BIGINT NOT NULL,
    requested_amount DECIMAL(14,2) NOT NULL,
    credit_score INT NOT NULL,
    risk_level ENUM('LOW','MEDIUM','HIGH','CRITICAL') NOT NULL,
    status ENUM('PENDING','APPROVED','REJECTED','UNDER REVIEW')
NOT NULL DEFAULT 'PENDING',
    interest_rate DECIMAL(5,2) NOT NULL,
    application_date DATE NOT NULL DEFAULT (CURRENT_DATE),
    FOREIGN KEY (client_id) REFERENCES clients(id),
    FOREIGN KEY (product_id) REFERENCES loan_products(id)
);

mvn flyway:migrate      # Appliquer
mvn flyway:info          # Statut
mvn flyway:clean         # DROP ALL (dangereux)

```

## 7.3 – Comparaison : Alembic vs Flyway

Aspect	Alembic	Flyway
Autogenerate	Oui (compare modèles vs DB)	Non (SQL écrit à la main)
Format migration	Python (opérations op.)	SQL pur
Nommage	Hash + message	V{n}__{desc}.sql strict
Rollback	downgrade() custom	Payant (Community: forward only)
Table de tracking	alembic_version	flyway_schema_history

<b>Aspect</b>	<b>Alembic</b>	<b>Flyway</b>
Intégration ORM	SQLAlchemy natif	Indépendant de l'ORM

---

## 8. Environnement Docker

```
# docker-compose.yml
version: "3.9"

services:
  mysql:
    image: mysql:8.0
    environment:
      MYSQL_ROOT_PASSWORD: root
      MYSQL_DATABASE: credit_scoring_db
      MYSQL_USER: analyst
      MYSQL_PASSWORD: analyst
    ports:
      - "3306:3306"
    volumes:
      - mysql_data:/var/lib/mysql
    healthcheck:
      test: ["CMD", "mysqladmin", "ping", "-h", "localhost"]
      interval: 10s
      retries: 5

  adminer:
    image: adminer:latest
    ports:
      - "8080:8080"
    depends_on:
      - mysql

  jupyter:
    build:
      context: .
      dockerfile: Dockerfile.jupyter
    ports:
      - "8888:8888"
    volumes:
      - ./notebooks:/home/jovyan/work
      - ./src:/home/jovyan/src
    depends_on:
      mysql:
        condition: service_healthy
    environment:
      - DATABASE_URL=mysql+pymysql://analyst:analyst@mysql:3306/
credit_scoring_db
```

```

volumes:
  mysql_data:

# Dockerfile.jupyter – Python + Java kernels
FROM jupyter/minimal-notebook:latest

USER root

# Java 21 + Maven
RUN apt-get update && \
    apt-get install -y openjdk-21-jdk maven && \
    apt-get clean

# IJava kernel (Java dans Jupyter)
RUN curl -L https://github.com/SpencerPark/IJava/releases/
download/v1.3.0/ijava-1.3.0.zip \
    -o /tmp/ijava.zip && \
    unzip /tmp/ijava.zip -d /tmp/ijava && \
    cd /tmp/ijava && python3 install.py --sys-prefix && \
    rm -rf /tmp/ijava*

USER $NB_UID

# Python deps via uv
COPY pyproject.toml /home/jovyan/
RUN pip install uv && \
    cd /home/jovyan && \
    uv pip install --system -r pyproject.toml

# pyproject.toml
[project]
name = "credit-scoring-analyst"
version = "0.1.0"
requires-python = ">=3.12"
dependencies = [
    "sqlalchemy[asyncio]>=2.0",
    "pymysql",
    "asyncmy",
    "pandas>=2.2",
    "alembic",
    "jupyterlab",
    "matplotlib",
    "cryptography",
]
<!-- pom.xml -->
<project>
  <modelVersion>4.0.0</modelVersion>

```

```
<groupId>com.creditscoring</groupId>
<artifactId>credit-scoring-analyst</artifactId>
<version>1.0-SNAPSHOT</version>

<properties>
    <java.version>21</java.version>
    <dflib.version>1.0.0-M22</dflib.version>
</properties>

<dependencies>
    <!-- DLib Core + JDBC + CSV -->
    <dependency>
        <groupId>org.dflib</groupId>
        <artifactId>dflib</artifactId>
        <version>${dflib.version}</version>
    </dependency>
    <dependency>
        <groupId>org.dflib</groupId>
        <artifactId>dflib-jdbc</artifactId>
        <version>${dflib.version}</version>
    </dependency>
    <dependency>
        <groupId>org.dflib</groupId>
        <artifactId>dflib-csv</artifactId>
        <version>${dflib.version}</version>
    </dependency>
    <!-- MySQL JDBC Driver -->
    <dependency>
        <groupId>com.mysql</groupId>
        <artifactId>mysql-connector-j</artifactId>
        <version>8.3.0</version>
    </dependency>

    <!-- JPA + Hibernate -->
    <dependency>
        <groupId>jakarta.persistence</groupId>
        <artifactId>jakarta.persistence-api</artifactId>
        <version>3.1.0</version>
    </dependency>
    <dependency>
        <groupId>org.hibernate.orm</groupId>
        <artifactId>hibernate-core</artifactId>
        <version>6.4.0.Final</version>
    </dependency>

    <!-- Flyway -->
    <dependency>
```

```

<groupId>org.flywaydb</groupId>
<artifactId>flyway-core</artifactId>
<version>10.0.0</version>
</dependency>
<dependency>
    <groupId>org.flywaydb</groupId>
    <artifactId>flyway-mysql</artifactId>
    <version>10.0.0</version>
</dependency>
</dependencies>
</project>

```

---

## Récapitulatif — Ce qu'il faut retenir

En tant que data analyst bancaire venant de Python :

1. **Ton workflow ne change pas** : SQL → DataFrame → Analyse → Export. Le domaine change (credit scoring au lieu d'académique), mais le pattern reste identique.
2. **SQLAlchemy 2.0** gère nativement les **Enum Python** (RiskLevel, ApplicationStatus) qui se mappent vers des ENUM MySQL. Les colonnes DECIMAL(12,2) pour les montants financiers sont typées avec Mapped[Decimal].
3. **Les requêtes bancaires** utilisent beaucoup de CASE WHEN (taux d'approbation), de subqueries (anomalies vs moyenne globale), et de GROUP BY multidimensionnels (risque × produit). SQLAlchemy Core exprime tout cela de manière composable avec case(), scalar\_subquery(), et func.
4. **DFlib** ne supporte pas nativement les opérations de binning (pd.cut) ni les lambda custom dans agg(). Pour ces cas, soit on fait le calcul côté SQL, soit on post-traite le DataFrame avec des filtres conditionnels chaînés.
5. **JPA** gère les enums avec @Enumerated(EnumType.STRING) — toujours utiliser STRING, jamais ORDINAL (qui stocke l'index et casse si on réordonne l'enum).
6. **Flyway** : les migrations contiennent les ENUM() MySQL en dur dans le DDL. Si on ajoute une valeur d'enum, il faut une nouvelle migration ALTER TABLE ... MODIFY COLUMN.